

# Foundations of Statistical and Machine Learning for Actuaries

## Resampling, cross-validation and regularisation


Edward (Jed) Frees, University of Wisconsin - Madison  
Andres M. Villegas, University of New South Wales

July 2025


# What is statistical (machine) learning?



# What is statistical (machine) learning?



# What is statistical (machine) learning?



## Prediction

- Predict outcomes of  $Y$  given  $X$ 
  - What it means isn't as important, it just needs accurate predictions
- Models tend to be more complex

## Inference

- Understand how  $Y$  is affected by  $X$
- Which predictors do we add? How are they related?
- Models tend to be simpler

# Regression vs. classification

---

## Regression

- $Y$  is quantitative, continuous
- Examples: Sales prediction, claim size prediction, stock price modelling

## Classification

- $Y$  is qualitative, discrete
- Examples: Fraud detection, face recognition, accident occurrence, death

More formally in regression we assume

---

$$Y = f(X) + \epsilon$$

- $Y$  is the outcomes, response, target variable
- $X := (X_1, X_2, \dots, X_p)$  are the features, inputs, predictors
- $\epsilon$  captures measurement error and other discrepancies

Our objective is to **find** an **appropriate**  $f$  for the problem at hand

# How to estimate $f$ ?

## Parametrics

- Make an assumption about the shape of  $f$
- Problem reduced down to estimating a few parameters
  - Works fine with limited data, provided assumption is reasonable
- Assumption strong: tends to miss some signal

## Non-parametric

- Make no assumption about  $f$ 's shape
- Involves estimating a lot of “parameters”
- Need lots of data
- Assumption weak: tends to incorporate some noise
- Be particularly careful re the risk of overfitting

# Parametrics example: Linear regression

Approximately a linear relationship between  $X$  and  $Y$

$$f(X) = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_p X_p$$

- The model is specified in terms of  $p + 1$  parameters  $\beta_0, \beta_1, \dots, \beta_p$ 
  - Use (training) data to produce estimates  $\hat{\beta}_0, \hat{\beta}_1, \dots, \hat{\beta}_p$
  - **Almost never correct**, but serves as a good and interpretable approximation.



# Non-parametrics example: K-nearest neighbours

---

KNN is one of the simplest non-parametric approaches

$$\hat{f}(x_0) = \frac{1}{K} \sum_{x_i \in \mathcal{N}_0} y_i$$

- can be pretty good for small  $p$  and large data sets (big  $N$ )
- need to choose the size of the value of  $K$ 
  - we will discuss other smoother versions such as local linear regression and splines in session 2

# Non-parametrics example: K-nearest neighbours

---

KNN is one of the simplest non-parametric approaches

$$\hat{f}(x_0) = \frac{1}{K} \sum_{x_i \in \mathcal{N}_0} y_i$$

- can be pretty good for small  $p$  and large data sets (big  $N$ )
- need to choose the size of the value of  $K$ 
  - we will discuss other smoother versions such as local linear regression and splines in session 2




# Non-parametrics example: K-nearest neighbours

KNN is one of the simplest non-parametric approaches

$$\hat{f}(x_0) = \frac{1}{K} \sum_{x_i \in \mathcal{N}_0} y_i$$

- can be pretty good for small  $p$  and large data sets (big  $N$ )
- need to choose the size of the value of  $K$ 
  - we will discuss other smoother versions such as local linear regression and splines in session 2

# How to choose $f$ ?



How do we decide which is the best model?

# Assessing model accuracy

---

We fit the model  $\hat{f}(x)$  to some **training** data  $Tr = \{x_i, y_i\}_{i=1}^n$ .

- We can compute the Training Mean Squared Error

$$MSE_{Tr} = \frac{1}{n} \sum_{i \in Tr} (y_i - \hat{f}(x_i))^2$$

# Assessing model accuracy

---

We fit the model  $\hat{f}(x)$  to some **training** data  $Tr = \{x_i, y_i\}_{i=1}^n$ .

- We can compute the Training Mean Squared Error

$$MSE_{Tr} = \frac{1}{n} \sum_{i \in Tr} (y_i - \hat{f}(x_i))^2$$

This tends to be biased to more overfit models!

# Assessing model accuracy

We fit the model  $\hat{f}(x)$  to some **training** data  $Tr = \{x_i, y_i\}_{i=1}^n$ .

- We can compute the Training Mean Squared Error

$$MSE_{Tr} = \frac{1}{n} \sum_{i \in Tr} (y_i - \hat{f}(x_i))^2$$

This tends to be biased to more overfit models!

We should instead use some fresh **test** data

$$Te = \{x_i, y_i\}_{i=1}^m.$$

- 

$$MSE_{Te} = \frac{1}{m} \sum_{i \in Te} (y_i - \hat{f}(x_i))^2$$

# Assessing model accuracy

---



# How do we calculate the test error?

---

1. The best solution is to use a large designated test set
  - Often not available
2. Make a mathematical adjustment to the training error rate
  - e.g. Cp statistic, AIC and BIC
3. Fit the model to a subset of the training observations
  - Use the remaining training observations as the test set

# k-fold Cross-validation

---

- Randomly divided the set of observations into  $K$  groups, or folds of approximately equal size
- the  $k^{\text{th}}$  fold is treated as a validation set
- the remaining  $K - 1$  folds make up the training set
- Repeat  $K$  times resulting  $K$  estimates of the test error

$$\text{CV}_{(K)} = \frac{1}{K} \sum_{k=1}^K \text{MSE}_k$$

- In practice  $K = 5$  or  $K = 10$

# k-fold Cross-validation

---

# Summary of key concepts

---

We have discussed key concepts in statistical/machine Learning

- Supervised learning vs. Unsupervised Learning
- Prediction vs. Inference
- Regression vs. Classification
- Parametric Vs. Non-Parametric
- Training MSE vs. Test MSE
- Cross-Validation

# Supervised learning: regression

# Regression vs. classification

---

## Regression

- $Y$  is quantitative, continuous
- Examples: Sales prediction, claim size prediction, stock price modelling

## Classification

- $Y$  is qualitative, discrete
- Examples: Fraud detection, face recognition, accident occurrence, death

More formally in regression we assume

---

$$Y = f(X) + \epsilon$$

- $Y$  is the outcomes, response, target variable
- $X := (X_1, X_2, \dots, X_p)$  are the features, inputs, predictors
- $\epsilon$  captures measurement error and other discrepancies

Our objective is to **find** an **appropriate**  $f$  for the problem at hand

# Can we predict house prices?



Source: <http://www.abc.net.au/news/2018-03-17/how-to-win-at-house-auction/9547166>

Output ( $Y$ ):

- House price

Input ( $X$ ):

- Home area
- Land area
- # of bedrooms
- # of bathrooms
- Neighbourhood
- Year built
- ...




# House Sales in King County, USA


Dataset from Kaggle of 21613 homes sold between May 2014 and May 2015.  
(<https://www.kaggle.com/harlfoxem/housesalesprediction/home>)

- price: Price is prediction target
- bedrooms: Number of Bedrooms
- bathrooms: Number of bathrooms/bedrooms
- sqft\_living: square footage of the home
- sqft\_lot: square footage of the lot
- floors: Total floors (levels) in house
- yr\_built: Built Year
- yr\_renovated: Year when house was renovated
- waterfront: House which has a view to a waterfront
- sqft\_above: square footage of house apart from basement

# House Sales in King County, USA



# House Sales in King County, USA



# Simple linear regression

- Approximately a linear relationship between  $X$  and  $Y$

$$Y = \beta_0 + \beta_1 X + \epsilon$$

- Use (training) data to produce estimates  $\hat{\beta}_0$  and  $\hat{\beta}_1$ 
  - Make predictions given  $X = x$


$$\hat{y} = \hat{\beta}_0 + \hat{\beta}_1 x$$

- Minimise the residual sum of squares (RSS)

$$\text{RSS} = \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \sum_{i=1}^n (y_i - \hat{\beta}_0 - \hat{\beta}_1 x)^2$$

# Simple linear regression: House prices

$$\text{price} = \beta_0 + \beta_1 \times \text{sqft\_living}$$



	Variable	estimate	std.error	p.value
1	(Intercept)	-47116.08	4923.34	0.00
2	sqft_living	281.96	2.16	0.00

# Multiple linear regression

---

- Extend the simple linear regression model to accommodate multiple predictors

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_p X_p + \epsilon$$

- $\beta_j$ : the average effect on  $Y$  of a one unit increase in  $X_j$ , holding all other predictors fixed

# Multiple linear regression: House prices

	Variable	estimate	std.error	p.value
1	(Intercept)	6289259.59	156282.14	0.00
2	bedrooms	-67820.03	2534.30	0.00
3	bathrooms	67280.69	4247.65	0.00
4	sqft_living	281.71	5.22	0.00
5	sqft_lot	-0.29	0.04	0.00
6	floors	43248.82	4526.50	0.00
7	yr_built	-3221.70	80.97	0.00
8	yr_renovated	6.69	4.74	0.16
9	waterfront	740322.15	20947.07	0.00
10	sqft_above	19.19	5.30	0.00

# Shortcomings of linear regression

---

1. **Prediction accuracy:** the linear regression fit often does not predict well, especially when  $p$  (the number of predictors) is large
2. **Model Interpretability:** linear regression freely assigns a coefficient to each predictor variable. When  $p$  is large, we may sometimes seek, for the sake of interpretation, a smaller set of **important variables**
3. **Non-linearities:** linear assumption is almost **always an approximation** – sometimes bad.




# Generalisations of the Linear Model

---

We discuss methods that expand the scope of linear models and how they are fit:


- *Regularised fitting*: Ridge regression and lasso
- *Classification problems*: logistic regression
- *Interactions*: Tree-based methods, bagging, random forests and boosting (these also capture non-linearities)

# Motivation: Linear Regression House prices




$$y = \beta_0 + \beta_1 x$$

- Underfit
- High bias



$$y = \beta_0 + \beta_1 x + \beta_2 x^2$$


- Just right



$$y = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3 + \beta_4 x^4 + \beta_5 x^5$$


- Overfit
- High variance

# Motivation: Linear Regression House prices




$$y = \beta_0 + \beta_1 x$$

- Underfit
- High bias



$$y = \beta_0 + \beta_1 x + \beta_2 x^2$$

- Just right



$$y = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3 + \beta_4 x^4 + \beta_5 x^5$$


- Overfit
- High variance

**Overfitting:** We have too many features, the model may fit the training set well ( $RSS \approx 0$ ), but fail to generalise to new cases (predict prices of new example)

# Overfitting with many features

Not unique to polynomial regression but also if lots of inputs ( $p$  large)


- $x_1 =$  Home area
- $x_2 =$  Land area
- $x_3 =$  # of bedrooms
- $x_4 =$  # of bathrooms
- $x_5 =$  Neighbourhood
- $x_6 =$  Year built
- $x_7 =$  Average income in the neighbourhood
- $x_8 =$  Kitchen size
- $\vdots$
- $x_{100}$



# Addressing Overfitting

There are several several options

1. Reduce number of features/variable
  - Manually
  - Subset selection algorithm
2. Regularisation
  - Keep all the features, but reduce magnitude of parameters  $\beta_i$ 
    - Works well when we have a lot of features, each of which contributes a bit to predicting  $y$



# Addressing overfitting via regularisation

---

$$\text{Total cost} = \text{Measure of Fit} + \text{Measure of Magnitude of Coefficient}$$

# Addressing overfitting via regularisation

---

$$\text{Total cost} = \underbrace{\text{Measure of Fit}}_{\text{RSS}} + \text{Measure of Magnitude of Coefficient}$$

# Addressing overfitting via regularisation

---

$$\text{Total cost} = \underbrace{\text{Measure of Fit}}_{\text{RSS}} + \underbrace{\text{Measure of Magnitude of Coefficient}}_{\beta_1^2 + \beta_2^2 + \dots + \beta_p^2}$$



# Ridge Regression


Minimise on  $\beta$ :

$$\sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2 = RSS + \lambda \sum_{j=1}^p \beta_j^2$$

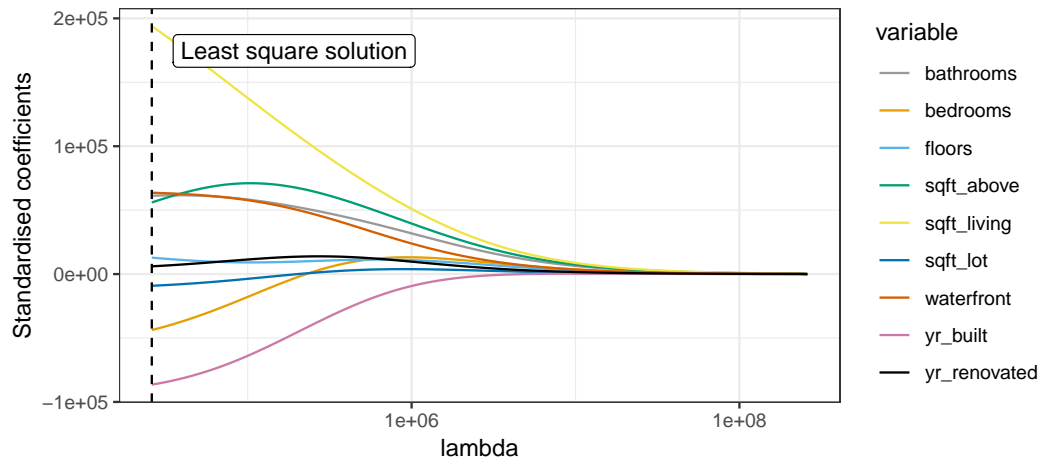
$\lambda$ : Tuning parameter = balance of fit and magnitude

- $\lambda \rightarrow \infty$ : Parameter estimates heavily penalised, coefficients pushed to zero, model is  $y_i = \hat{\beta}_0$
- $\lambda = 0$ : Parameter estimates not penalised at all, reduces to simple linear regression
  - obtain the best model which includes all parameters.


# Ridge Solutions paths: The house data



# Ridge Solutions paths: The house data



# Ridge Solutions paths: The house data




## Ridge Cross-Validation Solution: The house data


Variable	Estimate
(Intercept)	4461371.36
bedrooms	-23152.81
bathrooms	76655.13
sqft_living	155.79
sqft_lot	-0.11
floors	17043.10
yr_built	-2290.26
yr_renovated	27.11
waterfront	675263.65
sqft_above	85.58

Contains all variables so **still harder to interpret!**


# Thresholding ridge coefficients?



# Thresholding ridge coefficients?



# Thresholding ridge coefficients?





# Feature selection via regularisation

---

$$\text{Total cost} = \underbrace{\text{Measure of Fit}}_{\text{RSS}} + \underbrace{\text{Measure of Magnitude of Coefficient}}_{|\beta_1| + |\beta_2| + \dots + |\beta_p|}$$

# Lasso regression

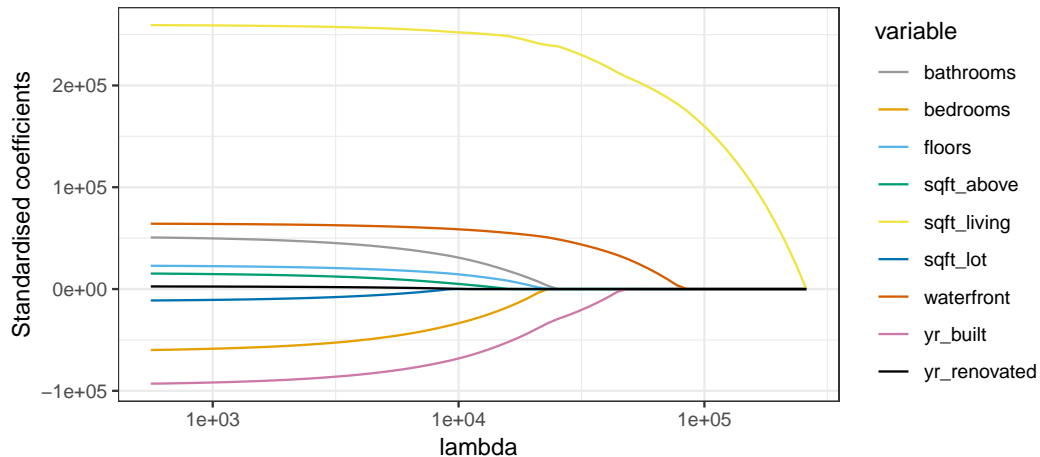
---

Minimise on  $\beta$ :


$$\sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p |\beta_j| = RSS + \lambda \sum_{j=1}^p |\beta_j|$$

- Only difference: penalties placed on absolute value of coefficient estimates
- Can force some of them to exactly zero: significantly easier to interpret model
- Has the effect of also performing some variable selection


# Lasso Solutions paths: The house data



# Lasso Solutions paths: The house data



# Lasso Solutions paths: The house data




## Lasso Cross-Validation Solution: The house data

Variable	estimate
(Intercept)	4166939.79
bedrooms	-30936.45
bathrooms	34095.80
sqft_living	272.38
sqft_lot	—
floors	22706.47
yr_built	-2134.77
yr_renovated	—
waterfront	659380.22
sqft_above	3.90

# Lasso vs. Ridge

## Lasso



## Ridge

